

A Practical Guide to Support Vector Classification

Chih-Jen Lin

Department of Computer Science
National Taiwan University




Talk at University of Freiburg, July 15, 2003

Motivation and Outline

- SVM: a hot machine learning issue
- However, many beginners get **unsatisfactory** accuracy at first
Some easy but significant steps missed
- This talk
 - Some **cookbook** approaches based on our experience serving users
 - No guarantee for the best accuracy but usually reasonable accuracy
 - Hope beginners get acceptable results fast and easily.
 - Challenging cases and further extensionWhat do we plan to add in LIBSVM

Basic Concepts of SVM



$$\min_{w, b, \xi} \quad \frac{1}{2} w^T w + C \sum_{i=1}^l \xi_i$$

subject to $y_i(w^T \phi(x_i) + b) \geq 1 - \xi_i, \quad \xi_i \geq 0, \quad i = 1, \dots, l.$

- Kernel: $K(x, y) = \phi(x)^T \phi(y)$

What Many Beginners are Doing Now

- Transfer data to the format of an SVM software
- May not conduct **scaling**
- **Randomly** try few parameters and kernels **without validation**
- Default parameters are surprisingly important
- If most users doing so, accuracy may not be satisfactory

Examples

	training data	testing data	features	classes	Accuracy by users	Accuracy by us
User 1	3,089	4,000	4	2	75.2%	96.9%
User 2	391	0	20	3	36%	85.2%
User 3	1,243	41	21	2	4.88%	87.8%

- User 1:

I am using libsvm in a astroparticle physics application .. First, let me congratulate you to a really easy to use and nice package.

Unfortunately, it gives me astonishingly bad results...

- Answer:

OK. Send me the data

- Answer:

I am able to get 97% test accuracy. Is that good enough for you ?

- User 1:

You earned a copy of my PhD thesis

- User 2:

I am a developer in a bioinformatics laboratory at ... We would like to use LIBSVM in a project ...

But results not good. 36% CV accuracy

- Answer:

OK. Send me the data

- Answer:

I am able to give 83.88% cv accuracy. Is that good enough for you ?

- User 2:
83.88% accuracy would be excellent...
- User 3:
I have problems getting the same result with SVM to compared to neural nets.
Right now I get a correct of 4.88%, which is **very bad** (neural net 70-90%).
- Answer
I play a bit your data. My testing accuracy is **87.8%**. Is this good for you ?
- User 3:
I found myself described in your talk ;-)

We Hope Users At Least Do

- The following procedure
 1. Conduct simple **scaling** on the data
 2. Consider **RBF** kernel $K(x, y) = e^{-\gamma \|x-y\|^2}$
 3. Use cross-validation to find the **best parameter** C and γ
 4. Use the best C and γ to **train the whole** training set
 5. Test

Why RBF

- Linear kernel: special case of RBF [Keerthi and Lin 2003]

- Polynomial: numerical difficulties

$$(< 1)^d \rightarrow 0, (> 1)^d \rightarrow \infty$$

More parameters than RBF

- tanh: still a **mystery**

May **not** be positive semi-definite

In [Lin and Lin 2003], for certain parameters, it **behaves like RBF**

Should avoid using tanh in general

Examples: Using the Proposed Procedure

User 1

- **Original** sets with default parameters

```
./svm-train train.1
```

```
./svm-predict test.1 train.1.model test.1.predict
```

→ Accuracy = 66.925%

- **Scaled** sets with default parameters

```
./svm-scale -s range1 train.1 > train.1.scale
```

```
./svm-scale -r range1 test.1 > test.1.scale
```

```
./svm-train train.1.scale
```

```
./svm-predict test.1.scale train.1.scale.model test.1.predict
```

→ Accuracy = 96.15%

- Scaled sets with **parameter selection**

```
$python grid.py train.1.scale
```

```
...
```

```
2.0 2.0 96.8922
```

(Best $C=2.0$, $\gamma=2.0$ with five-fold cross-validation
rate=96.8922%)

```
$/svm-train -c 2 -g 2 train.1.scale
```

```
$/svm-predict test.1.scale train.1.scale.model test.1.predict
```

→ Accuracy = 96.875%

User 2

- **Original** sets with default parameters

```
$/svm-train -v 5 train.2
```

→ Cross Validation Accuracy = 56.5217%

- **Scaled** sets with default parameters

```
$/svm-scale train.2 > train.2.scale
```

```
$/svm-train -v 5 train.2.scale
```

→ Cross Validation Accuracy = 78.5166%

- Scaled sets with **parameter selection**

```
$python grid.py train.2.scale
```

```
...
```

```
2.0 0.5 85.1662
```

→ Cross Validation Accuracy = 85.1662%

(Best $C=2.0$, $\gamma=0.5$ with five fold cross-validation rate=85.1662%)

User 3

- **Original** sets with default parameters

```
$/svm-train train.3
```

```
$/svm-predict test.3 train.3.model test.3.predict
```

→ Accuracy = 2.43902%

- **Scaled** sets with default parameters

```
./svm-scale -s range3 train.3 > train.3.scale
./svm-scale -r range3 test.3 > test.3.scale
./svm-train train.3.scale
./svm-predict test.3.scale train.3.scale.model test.3.predict
→ Accuracy = 12.1951%
```

- Scaled sets with **parameter selection**

```
$python grid.py train.3.scale
```

```
...
```

```
128.0 0.125 84.8753
```

(Best $C=128.0$, $\gamma=0.125$ with five-fold cross-validation rate=84.8753%)

```
./svm-train -c 128 -g 0.125 train.3.scale
```

```
./svm-predict test.3.scale train.3.scale.model test.3.predict
→ Accuracy = 87.8049%
```

Scaling

- Important for Neural Networks (Part 2 of NN FAQ)

Most reasons apply here

- Attributes in **greater numeric ranges may dominate**

$$K(x, y) = e^{-\gamma \|x - y\|^2}$$

- Simple linearly scaling each attribute to $[-1, +1]$ or $[0, 1]$.
- The **same** scaling factor for **testing**

Model Selection


- In fact, two-parameter search: C and γ
- We recommend a simple grid search using cross-validation
E.g. 5-fold CV on $C = 2^{-5}, 2^{-3}, \dots, 2^{15}$, $\gamma = 2^{-15}, 2^{-13}, \dots, 2^3$
- Why not more efficient methods

$$\text{leave-one-out error} \leq f(C, \gamma)$$


so

$$\min_{C, \gamma} f(C, \gamma)$$

- A path may be found



- Reasons for not using bounds (if two parameters)
 - Implementation more complicated
 - Psychologically, not feel safe
 - In practice: IJCNN competition:
97.09% and 97.83% using Radius Margin bounds for L1 and L2-SVM
98.59% using 25-point grid
2668, 1990, and 1293 testing errors
 - Bounds are useful if more than two parameters



- We propose that users do
 - Start from a **loose** grid
 - Identify good regions and use a **finer** grid
- The grid search tool in libsvm
- Easy parallelization

Every problem is independent

loos bounds: 20 steps \Rightarrow about 10×10 grids with five computers

Automatic load balancing

Example: Automatic Script

- User 1

```
$python easy.py train.1 test.1
```

```
Scaling training data...
```

```
Cross validation...
```

```
Best c=2.0, g=2.0
```

```
Training...
```

```
Scaling testing data...
```

```
Testing...
```

```
Accuracy = 96.875% (3875/4000) (classification)
```

- User 3

```
$python easy.py train.3 test.3
```

```
Scaling training data...
```

```
Cross validation...
```

Best $c=128.0$, $g=0.125$

Training...

Scaling testing data...

Testing...

Accuracy = 87.8049% (36/41) (classification)

Challenges

- Is the procedure good enough ?
Good for **some median-sized** data sets
- Difficult problems: this procedure **not enough**
 - Too much training time
 - Low accuracy
- Extension of the procedure ?
- What are we going to include in LIBSVM?

Feature Selection

- Too many (non-zero) features

Examples here: 4, 20, 21 features \ll #data

- RBF kernel

$$K(x, y) = e^{-\gamma \|x - y\|^2}$$

Irrelevant attributes cause problems

- How about

$$K(x, y) = e^{-\sum_{i=1}^n \gamma_i (x_i - y_i)^2}$$

Difficult to choose γ_i

Possible approaches (e.g. [Chapelle et al. 2002]):

$$\text{leave-one-out error} \leq f(C, \gamma_1, \dots, \gamma_n)$$

A non-convex problem. Difficult and unstable.

- Feature selection before training SVM

SVM can help feature selection as well

E.g. linear SVM

$$f(x) = w^T x + b$$

Choose indices with large $|w_i|$ [Guyon et al. 2002]

- Overall, a very difficult issue

Not sure if a **simple and systematic** procedure available ?

Probability Estimates

- SVM outputs **decision values only**
- Probability estimates for **two-class** SVM:
 - Platt's sigmoid approximation
 - Isotonic regression
 - SVM density estimation ?

We are conducting a serious evaluation

- Multi-class probability estimate
Related to multi-class classification

Currently LIBSVM uses 1vs1 (after an evaluation in [Hsu and Lin, 2002])

10 classes: 45 SVMs, 0vs1, 0vs2, ..., 8vs9

Given $r_{ij} \approx P(y = i \mid y = i \text{ or } j)$, estimate $P(y = i)$

An issue for all binary classification methods

New and stable methods proposed in [Wu et al., 2003]

- All these are about ready

The main addition to next version of LIBSVM

Unbalanced Data

- Many information retrieval users ask about ROC curve and adjusting precision/recall

Not accuracy any more

- Three ways to generate ROC curves


- Adjust b of

$$f(x) = w^T x + b$$

- Unbalanced cost function

$$\min_{w,b,\xi} \frac{1}{2} w^T w + C_+ \sum_{i:y_i=1} \xi_i + C_- \sum_{i:y_i=-1} \xi_i$$

- Rank by probability output + cross validation (now available)
- Which one is more useful ?



- Goal: an **integrated** tool so users can easily adjust cost matrices or the relation of TP, TN, FT, FN

Conclusions

- Still a long way to serve all users' needs but we are trying
- We hope more users can benefit from this research and eventually SVM can be an **easy-to-use** classification method
- Slides based on
Chih-Wei Hsu, Chih-Chung Chang, and Chih-Jen Lin
A Practical Guide to Support Vector Classification <http://www.csie.ntu.edu.tw/~cjlin/papers/guide/guide.pdf>
- LIBSVM available at
<http://www.csie.ntu.edu.tw/~cjlin/libsvm>
- We thank all users for their comments